

Team Brief — SIH26152 Internal Round

For the teammates building the presentation. Everything the deck needs: what the problem statement asks, what is actually running, what deliberately is not, and the numbers to put on slides.

Styled version (same content, nicer to read): <https://claude.ai/code/artifact/e2778199-b136-47ba-8987-a8b404443f96> If you edit one, edit the other — a stale copy is worse than one copy.

Problem statement	SIH26152 — Social Media Analytics · Sponsor NTRO
Internal round	29–30 Aug 2026
After that	Idea submission 20 Sept · Grand Finale Dec 2026 (36h)
State as of	b33f67d
Repo	Krishna-Das20/SIH-26152

Each section below maps to roughly one slide. **Every number here was read off the running system on 26 Aug 2026 — none are aspirational.**

If a number is not in this document, do not put it on a slide. That rule is the whole reason the project reads as credible.

1. What the problem statement actually asks for

Slide 1–2 · Problem framing

NTRO wants a system that reads public social media and tells an analyst not just *what* is being said, but *who* is saying it, *how they feel*, and *how it spreads*. Five named components:

#	Component	What it demands
A	Continuous data collection	Tiered: Essential — X & Telegram. Desirable — Instagram & Facebook. Appreciable — Reddit or YouTube.
B	Multi-dimensional sentiment	Beyond positive/negative — sarcasm, anxiety, excitement, supportive, against.
C	Demographic profiling	Automated inference of age, location, language, interests.
D	Trend & topic detection	Identify, rank, and predict emerging topics in real time.
E	Link analysis	Network topology — communities, influencers, how content travels.

The line that decides the winner

The statement says plainly:

"Combining these four vectors using AI is the key to unlocking true audience intelligence."

Most teams will build five separate dashboards and stop. The brief is that the *combination* is the product. Section 3 is our answer to that sentence, and it should get the most stage time.

2. What is built, stated honestly

Slide 3 · Component status matrix

#	Component	Status	What is actually behind it
A	Ingestion	4 of 6 live	Six connectors written against real APIs. YouTube, Telegram, Instagram live; Facebook token valid but feed scope missing; X needs funding; Reddit is policy-gated.
B	Sentiment & emotion	Real models	Five transformers on a Python service — RoBERTa sentiment, GoEmotions, a sarcasm classifier, toxicity, MiniLM embeddings. 352 of 352 posts transformer-scored.
C	Demographics	🌀 Rules only	Regex and lexicons, no ML. Reports its own coverage and returns null rather than guessing. Our known weak point — own it, don't hide it.
D	Trends	No forecast	Real z-score spike detection over 2,035 time buckets. Ranks and detects; does not yet predict.
E	Link analysis	Real algorithms	Louvain community detection, Brandes betweenness, damped PageRank. Verified against hand-computed modularity.
+	Cross-vector brief	Differentiator	The fusion layer. Six finding types that only exist by intersecting vectors. See section 3.

+	Narrative mutation	New	Clusters posts by semantic similarity, scores how a narrative shifts between its early and late halves.
---	--------------------	-----	---

The rule that makes all of this defensible

Nothing in the system fabricates a number. Where a platform does not report a follower count, the field is `null` and the UI says "Unknown". Where the ML service is unreachable, the page says so instead of silently degrading. Seven places that previously generated `Math.random()` values were removed.

Say this out loud to the judges. Half the projects in the room will have plausible-looking numbers that came from nowhere, and a judge who probes one will find it. Ours survives the probe.

3. The differentiator — findings no single vector can see

Slide 4–5 · Spend the most time here

Every other panel answers a question you already knew to ask. The Audience Intelligence Brief answers questions an analyst didn't know to ask, by intersecting vectors that are normally reported separately.

Three real findings from the current corpus, verbatim from the running system:

[HIGH] Negative sentiment is concentrated in one community, not spread evenly `sentiment × network`

Community 4 (9 accounts, topic `#IndiaSemiconductorMission`) is 56% negative, against 0% in Community 5. A platform-wide average would have hidden this 56-point split entirely.

[NOTABLE] Sentiment on #GovernmentBudgetEconomy turned negative mid-conversation `trends × sentiment`

Average sentiment moved from 0.33 to -0.20 across 9 posts, turning around 17 Mar 2026. Tracking the topic total alone would show only that it was busy.

[NOTABLE] 7 posts read as positive but are sarcastic `sentiment × sarcasm`

A polarity-only model would score these as support. The sarcasm classifier flags them, and the brief reports the gap between apparent and actual stance.

Every finding carries the evidence it was derived from, so a reader can check it rather than trust it. There is a **Show evidence** toggle in the UI — open it on stage.

The one-sentence pitch

"A dashboard tells you sentiment is 62% negative. Ours tells you the negativity lives in one 9-account community, is about semiconductors, and is being carried between clusters by one specific broker — which is the difference between a chart and an intelligence product."

4. Numbers you can safely put on a slide

Slide 6 · Evidence

352	posts in the demo corpus — real, captured live
352	transformer-scored — 100%, no lexicon fallback
298	accounts in the graph, 174 real edges
5	transformer models running locally on CPU
7	cross-vector findings, each with evidence
3	platforms with live data — 6 connectors written

Corpus composition

Platform	Posts	Tier in the PS
Instagram	160	Desirable
YouTube	152	Appreciable
Telegram	40	Essential

Emotion distribution — component B in one table

This is the table that proves "multi-dimensional" is real and not a relabelled positive/negative axis. Note that *supportive* and *against* are **stances**, not moods — they are derived separately from emotional valence, because "I'm nervous about this" is worry, not opposition.

Dimension	Posts		Dimension	Posts
Supportive	128		Sadness	7
Neutral	87		Anger	4
Excitement	66		Fear	1
Joy	51		Anxiety	0
Against	8		Sarcastic	3%

Stance split: **193 supportive · 12 opposing · 147 neutral**.

Anxiety at zero is a *true* reading of this corpus — it is mostly fan comments and product announcements, not a crisis. **Do not treat a zero as a bug.**

Network

Measure	Value	Note
Nodes / edges	298 / 174	No synthesised edges — every link is a real reply or mention
Communities	5	Excluding 119 isolated single-account groups

Modularity Q	0.2796	Below 0.3 — see section 7
Top influencers ranked	8	By PageRank and betweenness, not follower count

5. Architecture

Slide 7 · One diagram

Two processes, deliberately. The dashboard is stateless and deploys anywhere; the model service holds several gigabytes of weights in memory and cannot.

Layer	Technology	Responsibility
Dashboard	Next.js 14 · TypeScript · Tailwind	UI, API routes, graph rendering, all analytics maths
Model service	Python · FastAPI · PyTorch	Sentiment, emotion, sarcasm, toxicity, embeddings
Storage	MongoDB Atlas	Post persistence, deduplication, encrypted user tokens
Ingestion	Six connectors, one interface	Fault-isolated — one platform failing never blocks the others

The models

Purpose	Model	Disk
Sentiment	<code>cardiffnlp/twitter-roberta-base-sentiment-latest</code>	957 MB
Emotion	<code>SamLowe/roberta-base-go_emotions</code>	479 MB
Sarcasm	<code>hallisky/sarcasm-classifier-gpt4-data</code>	1.4 GB
Toxicity	<code>unitary/toxic-bert</code>	418 MB
Embeddings	<code>sentence-transformers/all-MiniLM-L6-v2</code>	88 MB

The sarcasm model was chosen **empirically, not from a blog post** — four candidates were tested and this one scored 0.999 on clear sarcasm with zero false positives on the control set. That anecdote plays well if a judge asks how models were selected.

Graceful degradation

If the model service is down, every panel except Narratives keeps working, because the demo corpus ships pre-scored. Worth demonstrating if there's time — **stop the Python service on stage and reload**. Nothing breaks; the one affected panel names the cause.

6. Platform access is the real constraint

Slide 8 · The slide that shows we did the homework

Every connector is written and calls the real API. What separates live from dormant is credentials and platform policy — **not code**. We verified each of these ourselves rather than repeating what blogs claim:

Platform	Status	What we found
Telegram	Live	The only platform with an open route left. Public channel previews need no credentials at all.
YouTube	Live	Data API v3, 10,000 units/day free. Search costs 100 units, comments cost 1.
Instagram	Live	Graph API with a permanent Page token. Basic Display API retired Dec 2024; unauthenticated routes return 429.
Facebook	Token valid	Reads Page metadata; the feed needs one more permission scope we haven't granted.
X (Twitter)	🌀 Needs funding	Pay-per-call since Feb 2026 — no free tier. Roughly \$24/month at our volume. Code is written and waiting.
Reddit	🌀 Policy-gated	Responsible Builder Policy routes new Data API apps through a request process aimed at moderation use cases. Public JSON returns 403.

Why this is a strength, not an excuse

Teams that claim six live platforms are either paying for X or scraping in breach of terms. Naming the exact blocker for each platform, with the evidence, demonstrates engineering judgement.

The **essential** tier in the problem statement is X *and* Telegram — Telegram is live, and X is a **purchase decision, not an engineering one**.

7. What we deliberately did not do

Slide 9 · Put this on a slide — judges reward it

Volunteering the gaps before a judge finds them converts a weakness into evidence of rigour.

Gap	Why it is open	Plan by December
Demographics has no ML	Regex and lexicons. Coverage is honest: 58% language, 6% age, 0% location.	Fine-tune a classifier; the interface already returns <code>null</code> , so nothing else changes.
Trends do not forecast	Component D asks for prediction. We detect and rank; we do not yet predict.	Add a time-series model over the existing 2,035-bucket history.
No accuracy evaluation	We report which engine scored each post, but have no precision/recall against a labelled set.	Hand-label 200 posts, publish a confusion matrix.
Modularity is 0.28	Instagram's 160 comments form one star around a single account, diluting the structure. Below the 0.3 "meaningful clustering" convention.	Ingest more discussion-heavy sources. This reached 0.83 on an earlier corpus.

⚠ Before the 29th The modularity number is the one gap that is cheap to close: adding two or three more YouTube queries on a single contested topic and re-freezing the corpus lifts it above 0.3 *and* gives the narrative tracker real discourse to work with. Worth doing before the deck is finalised so the number on the slide is the better one.

8. Demo script

Live segment · Target 5 minutes

Start both services **before you walk in** — the model service takes about 14 seconds to load weights and another 20 on its first request. Have the dashboard already open.

01 · Open the dashboard cold No login, no network dependency.

"352 real posts from three platforms, already analysed. Nothing here is sample data."

02 · Emotion radar — component B Point at Supportive and Against sitting beside Joy and Anger.

"Stance and mood are computed separately. Someone anxious about a policy is not opposing it, and most tools conflate those."

03 · Network graph — component E Show the communities.

"Real Louvain and Brandes betweenness. Every edge is an actual reply or mention — we removed the code that invented edges to make the graph look fuller."

04 · The Audience Intelligence Brief — the differentiator Read the top finding aloud, then hit **Show evidence**. *Spend the most time here.*

"This finding does not exist in any single panel. It required sentiment and network topology at the same time."

05 · Narratives page Show clustering by meaning rather than hashtag. Note that most scores read "insufficient data".

"That's the system refusing to score a two-post cluster. It would rather say nothing than say something unsupported."

06 · If time allows — kill the model service Reload. Everything still works; the one dependent panel names the cause.

"Degradation is visible, never silent."

⚠ Risk control Record a screen-capture of this exact run the night before and keep it on the presenting laptop. If the venue Wi-Fi or the laptop misbehaves, you present the video and lose nothing — the corpus is frozen, so the recording and the live run show identical numbers.

9. Deployment

Slide 10 · Path to production

The dashboard is already deployed

Next.js runs on Vercel today (<https://sih-26152.vercel.app>) and serves the same frozen corpus with identical analysis. That part is solved.

The model service cannot run on Vercel

This is not a configuration problem — it's a structural mismatch, and it's worth being able to explain why:

- The five models occupy **3.3 GB on disk** and the warmed service holds **3.1 GB resident**. Vercel functions cap at roughly 3 GB of memory and a far smaller deployment bundle — PyTorch alone exceeds the bundle limit before any weights are added.
- Serverless functions are ephemeral. Every cold start would reload the weights: about **14 seconds** to load plus **20 seconds** on the first request, against execution-time limits measured in seconds.

The correct architecture is exactly what we have: a stateless dashboard on Vercel calling a long-lived model service over HTTP, pointed at by one environment variable (`ML_API_URL`).

Measured requirements for sizing the host

Resource	Measured	Provision
Memory	3.1 GB resident	4 GB minimum, 8 GB comfortable
Disk	3.3 GB weights	~10 GB with the Python environment
CPU	9.3 posts/sec	2 vCPU is enough; no GPU needed
Cold start	~34 s to first result	Keep one instance warm
Embeddings	547 texts/sec	Negligible — never the bottleneck

Where to host it

Option	Spec	Cost	Verdict
Hugging Face Spaces	2 vCPU · 16 GB	Free	Start here. Fits comfortably, Docker support, public HTTPS, and the weights are already hosted there. Free tier sleeps when idle.
Oracle Cloud Free Tier	4 ARM cores · 24 GB	Free, always on	Best free option that never sleeps. More setup — ARM builds of PyTorch.
Azure	B2s VM · 2 vCPU · 4 GB	~\$30/mo	Azure for Students credit covers it.
AWS	EC2 t3.large · 8 GB	~\$60/mo	Works, but the free-tier t2.micro at 1 GB will not — do not promise it on a slide.
GCP Cloud Run	4 GB · scales to zero	Usage-based	Cold start makes the first request painful unless you pay for a warm instance.

Recommendation for the internal round Do not deploy the model service before the 29th. Demo from a laptop. The dashboard is already live on Vercel for anyone who wants a link, the corpus is frozen so results are identical either way, and a rushed deploy two days out is pure downside risk. Hugging Face Spaces is the September task, ahead of idea submission on the 20th.

10. Questions judges will ask

Preparation only — not a slide

Rehearse these. **The honest answer is stronger than the impressive one in every case below.**

Q · Is this real data or a mock-up? **Real.** 352 posts pulled from live APIs and frozen so the demo doesn't depend on venue Wi-Fi. Every one carries a transformer score, and we can show the ingestion running live if you'd like.

Q · Why only three platforms? Six connectors are written against real APIs. X charges per call since February and we haven't funded it; Reddit gates new API apps behind a policy aimed at moderation tools; Facebook needs one more permission scope. Those are procurement and policy blockers, not missing code — and we'd rather name them than scrape in breach of terms.

Q · How accurate is the sentiment analysis? **We don't have a number yet, and we won't invent one.** We use published models with known benchmarks and record which engine scored each post. Hand-labelling 200 posts for a confusion matrix is our next task.

Q · Isn't demographic inference just guessing? It would be if we let it. Ours is rules-based and reports its own coverage — 58% for language, 6% for age, 0% for location on this corpus. Unknown values stay unknown rather than defaulting to a plausible-looking value. Making it a real classifier is scheduled work.

Q · What stops this being another sentiment dashboard? The cross-vector brief. The problem statement says combining the vectors is the point, so we built the layer that does it — findings like "the negativity is concentrated in one nine-account community, about semiconductors, carried by one broker." No single panel can produce that.

Q · Could this scale to national volume? The analytics scale fine — the graph algorithms are standard and the ingestion is fault-isolated per platform. The model service is the constraint at 9 posts per second per instance on CPU; it's stateless, so it scales horizontally, and a GPU would move it by roughly an order of magnitude.

Q · What about user privacy? Only public content is collected. For the multi-tenant path, tokens are encrypted at rest with AES-256-GCM bound to the user and provider, and there's a data-deletion endpoint. For the demo, no personal accounts are connected at all.

Brief compiled 26 Aug 2026. All figures read from the running system. Related: [PROGRESS.md](../PROGRESS.md) for current status · [AGENTS.md](../AGENTS.md) for architecture.